

# Reconhecimento de imagens - Convolutional Neural Networks

## Diferença entre as espécies de pistachios, Siirt e Kirmizi

- Beatriz Carvalho e Carolina Aquino

Dataset original com 2149 imagens e reduzido para 998 imagens.

## Importar e reduzir dataset

```
In [5]: !pip install tensorflow
!pip install keras

Requirement already satisfied: tensorflow in ./opt/anaconda3/lib/python3.9/site-packages (2.11.0)
Requirement already satisfied: flatbuffers>=2.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (23.1.21)
Requirement already satisfied: packaging in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (21.3)
Requirement already satisfied: setuptools in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (61.2.0)
Requirement already satisfied: opt-einsum>=2.3.2 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (3.3.0)
Requirement already satisfied: wrapt>=1.11.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (1.12.1)
Requirement already satisfied: absl-py>=1.0.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (1.4.0)
Requirement already satisfied: tensorflow<2.12,>=2.11 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (2.11.2)
Requirement already satisfied: typing-extensions>=3.6.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (4.1.1)
Requirement already satisfied: protobuf<3.20,>=3.9.2 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (3.19.1)
Requirement already satisfied: termcolor>=1.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (2.2.0)
Requirement already satisfied: libclang<13.0.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (15.0.6.1)
Requirement already satisfied: astunparse>=1.6.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (1.6.3)
Requirement already satisfied: numpy>=1.20 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (1.21.0)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (1.42.0)
Requirement already satisfied: h5py>=2.9.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (3.6.0)
Requirement already satisfied: tensorflow-estimator<2.12,>=2.11.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (2.11.0)
Requirement already satisfied: google-pasta<0.1.4 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (0.2.0)
Requirement already satisfied: six>=1.12.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (1.16.0)
Requirement already satisfied: gast<0.4.0,>=0.2.1 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (0.4.0)
Requirement already satisfied: pyrsistent<3.0.5,>=2.0.2 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow) (0.30.0)
Requirement already satisfied: wheel<1.0,>=0.23.0 in ./opt/anaconda3/lib/python3.9/site-packages (from astunparse==1.6.0->tensorflow) (0.37.1)
Requirement already satisfied: markdown>=2.8.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow<2.12,>=2.11->tensorflow) (3.3.4)
Requirement already satisfied: tensorflow-data-server<0.7.0,>=0.6.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow<2.12,>=2.11->tensorflow) (0.6.1)
Requirement already satisfied: werkzeug>=1.0.1 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow<2.12,>=2.11->tensorflow) (2.0.3)
Requirement already satisfied: request>=3,>=2.21.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow<2.12,>=2.11->tensorflow) (2.27.1)
Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow<2.12,>=2.11->tensorflow) (0.4.6)
Requirement already satisfied: tensorflow-plugin-wlit>=1.6.0 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow<2.12,>=2.11->tensorflow) (1.8.1)
Requirement already satisfied: google-auth<3,>=1.6.3 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow<2.12,>=2.11->tensorflow) (1.33.0)
Requirement already satisfied: cachetools<5.0,>=2.0.0 in ./opt/anaconda3/lib/python3.9/site-packages (from google-auth<3,>=1.6.3->tensorflow<2.12,>=2.11->tensorflow) (4.2.2)
Requirement already satisfied: pyasn1-modules<0.2.1 in ./opt/anaconda3/lib/python3.9/site-packages (from google-auth<3,>=1.6.3->tensorflow<2.12,>=2.11->tensorflow) (0.2.8)
Requirement already satisfied: rsa<5,>=3.1.4 in ./opt/anaconda3/lib/python3.9/site-packages (from google-auth<3,>=1.6.3->tensorflow<2.12,>=2.11->tensorflow) (4.7.2)
Requirement already satisfied: requests-oauthlib<0.7.0 in ./opt/anaconda3/lib/python3.9/site-packages (from google-auth-oauthlib<0.5,>=0.4.1->tensorflow<2.12,>=2.11->tensorflow) (1.3.1)
Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in ./opt/anaconda3/lib/python3.9/site-packages (from pyasn1-modules==0.2.1->tensorflow<2.12,>=2.11->tensorflow) (0.4.8)
Requirement already satisfied: charset-normalizer<=2.0.0 in ./opt/anaconda3/lib/python3.9/site-packages (from requests<3,>=2.21.0->tensorflow<2.12,>=2.11->tensorflow) (2.0.4)
Requirement already satisfied: idna<4,>=2.5 in ./opt/anaconda3/lib/python3.9/site-packages (from requests<3,>=2.21.0->tensorflow<2.12,>=2.11->tensorflow) (3.3.2)
Requirement already satisfied: certifi<=2017.4.17 in ./opt/anaconda3/lib/python3.9/site-packages (from requests<3,>=2.21.0->tensorflow<2.12,>=2.11->tensorflow) (2021.10.8)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in ./opt/anaconda3/lib/python3.9/site-packages (from requests<3,>=2.21.0->tensorflow<2.12,>=2.11->tensorflow) (1.26.0)
Requirement already satisfied: oauthlib<3.0.0 in ./opt/anaconda3/lib/python3.9/site-packages (from requests-oauthlib<0.7.0->google-auth-oauthlib<0.5,>=0.4.1->tensorflow<2.12,>=2.11->tensorflow) (3.2.2)
Requirement already satisfied: tensorflow-io-gcs-filesystem<=0.23.1 in ./opt/anaconda3/lib/python3.9/site-packages (from tensorflow<2.12,>=2.11->tensorflow) (0.37.1)
Requirement already satisfied: keras in ./opt/anaconda3/lib/python3.9/site-packages (from packaging->tensorflow) (3.0.4)
```

## Construção do modelo com Keras

```
In [8]: from tensorflow import keras
from tensorflow.keras import layers
import keras

In [9]: inputs = keras.Input(shape=(224, 224, 3))
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(inputs)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.summary()

Model: "model"
Layer (type) Output Shape Param #
-----
input_1 (InputLayer) [(None, 224, 224, 3)] 0
conv2d (Conv2D) (None, 222, 222, 32) 896
max_pooling2d (MaxPooling2D) (None, 111, 111, 32) 0
conv2d_1 (Conv2D) (None, 109, 109, 64) 18496
max_pooling2d_1 (MaxPooling2D) (None, 54, 54, 64) 0
conv2d_2 (Conv2D) (None, 52, 52, 128) 73856
max_pooling2d_2 (MaxPooling2D) (None, 26, 26, 128) 0
conv2d_3 (Conv2D) (None, 24, 24, 256) 295168
max_pooling2d_3 (MaxPooling2D) (None, 12, 12, 256) 0
conv2d_4 (Conv2D) (None, 10, 10, 256) 590880
max_pooling2d_4 (MaxPooling2D) (None, 5, 5, 256) 0
flatten (Flatten) (None, 6400) 0
dense (Dense) (None, 1) 6401
-----
Total params: 984,897
Trainable params: 984,897
Non-trainable params: 0
```

## Tratamento das imagens

```
In [10]: from keras import optimizers
model.compile(loss='binary_crossentropy',
              optimizer='rmsprop',
              metrics=['accuracy'])

In [11]: from keras.preprocessing.image import ImageDataGenerator
train_data_gen = ImageDataGenerator(rescale=1./255)
test_data_gen = ImageDataGenerator(rescale=1./255)
train_generator = train_data_gen.flow_from_directory(
    new_base_dir / 'train',
    target_size=(224, 224),
    batch_size=200,
    class_mode='binary')

validation_generator = test_data_gen.flow_from_directory(
    new_base_dir / 'validation',
    target_size=(224, 224),
    batch_size=200,
    class_mode='binary')
```

Found 598 images belonging to 2 classes.

Found 400 images belonging to 2 classes.

## Ajuste do modelo

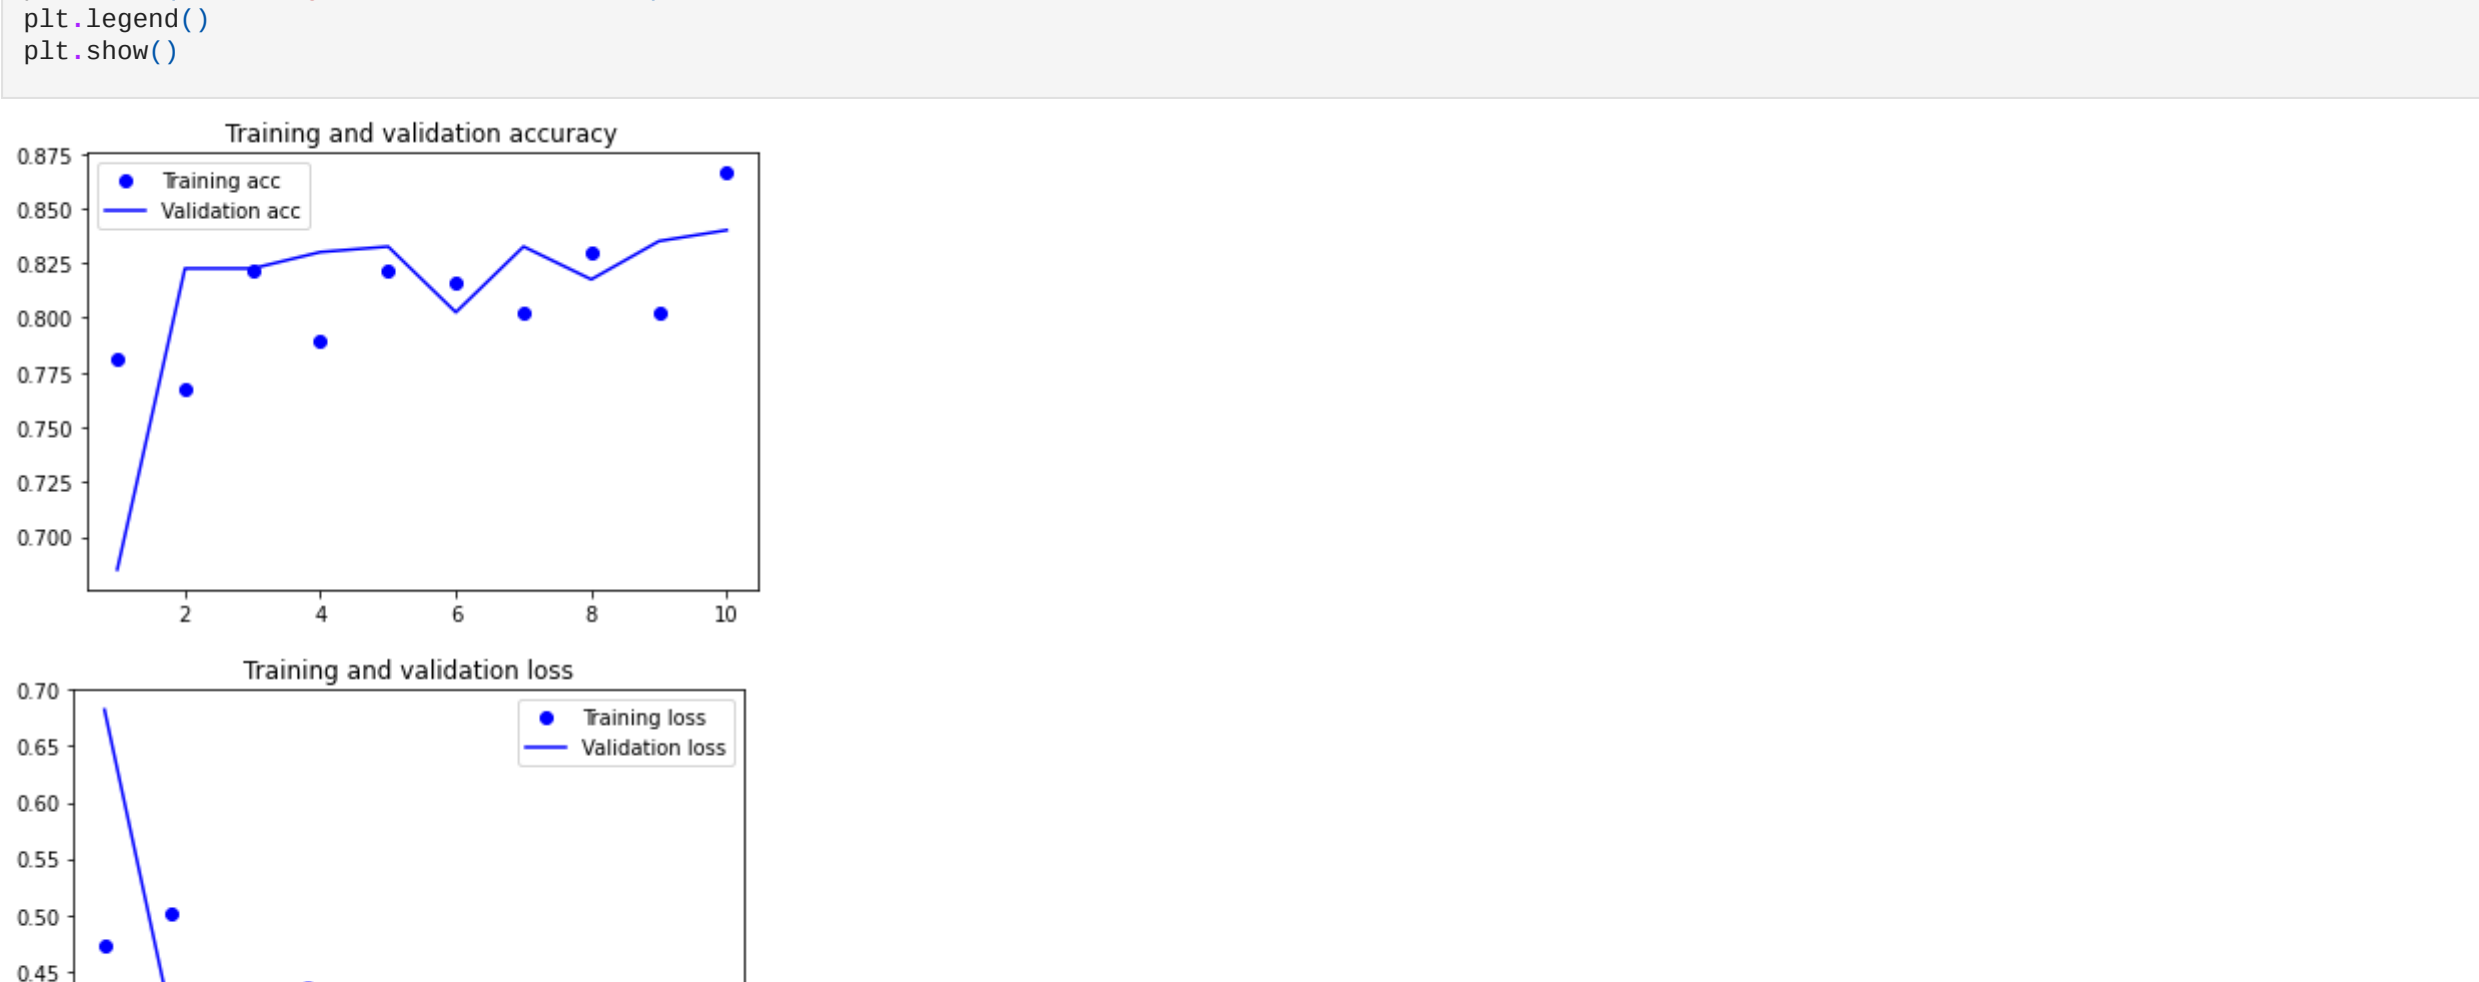
```
In [14]: callbacks = [
keras.callbacks.ModelCheckpoint(
    filepath='convnet_from_scratch.keras',
    save_best_only=True,
    monitor='val_loss')]

history = model.fit(
    train_generator,
    epochs=10,
    validation_data=validation_generator,
    callbacks=callbacks)

Epoch 1/10
3/3 [=====] - 30s 10s/step - loss: 0.4727 - accuracy: 0.7809 - val_loss: 0.6825 - val_accuracy: 0.6850
Epoch 2/10
3/3 [=====] - 30s 11s/step - loss: 0.5011 - accuracy: 0.7676 - val_loss: 0.4033 - val_accuracy: 0.8225
Epoch 3/10
3/3 [=====] - 29s 10s/step - loss: 0.3951 - accuracy: 0.8211 - val_loss: 0.4099 - val_accuracy: 0.8225
Epoch 4/10
3/3 [=====] - 29s 10s/step - loss: 0.4360 - accuracy: 0.7893 - val_loss: 0.3887 - val_accuracy: 0.8300
Epoch 5/10
3/3 [=====] - 29s 10s/step - loss: 0.3788 - accuracy: 0.8211 - val_loss: 0.3878 - val_accuracy: 0.8325
Epoch 6/10
3/3 [=====] - 28s 10s/step - loss: 0.3979 - accuracy: 0.8161 - val_loss: 0.4370 - val_accuracy: 0.8025
Epoch 7/10
3/3 [=====] - 28s 10s/step - loss: 0.4060 - accuracy: 0.8027 - val_loss: 0.3802 - val_accuracy: 0.8325
Epoch 8/10
3/3 [=====] - 28s 10s/step - loss: 0.3664 - accuracy: 0.8294 - val_loss: 0.4079 - val_accuracy: 0.8175
Epoch 9/10
3/3 [=====] - 30s 10s/step - loss: 0.4061 - accuracy: 0.8027 - val_loss: 0.3628 - val_accuracy: 0.8350
Epoch 10/10
3/3 [=====] - 29s 10s/step - loss: 0.3303 - accuracy: 0.8662 - val_loss: 0.3489 - val_accuracy: 0.8400
```

In [15]: model.save('pistachio\_small.h5')

```
In [16]: import matplotlib.pyplot as plt
```




## Melhorar o modelo - Data Augmentation e Droupout

```
In [59]: datagen = ImageDataGenerator(rotation_range=40,
width_shift_range=0.2,
height_shift_range=0.2,
shear_range=0.2,
zoom_range=0.2,
horizontal_flip=True,
fill_mode='nearest')

In [60]: from keras.preprocessing import image
names = [os.path.join(new_base_dir / 'train' / 'siirt', fname) for fname in os.listdir(new_base_dir / 'train' / 'siirt')]
img_path = names[3]
img = image.load_img(img_path, target_size=(224, 224))
```

```
In [61]: x = image.img_to_array(img)
x = x.reshape(1,1) * x.shape

i = 0
for batch in datagen.flow(x, batch_size=1):
    plt.figure(1)
    imgplot = plt.imshow(image.array_to_img(batch[0]))
    i += 1
    if i % 4 == 0:
        break
plt.show()
```



## Novo ajuste do modelo

```
In [62]: import tensorflow as tf
from tensorflow.keras import optimizers

model = keras.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
input_shape=(224, 224, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
model.compile(optimizer='keras.optimizers.RMSprop(
    learning_rate=0.001),
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

```
In [63]: callbacks = [
keras.callbacks.ModelCheckpoint(
    filepath='feature_extraction.keras',
    save_best_only=True,
    monitor='val_loss')
]
```

```
In [64]: train_data_gen = ImageDataGenerator(
rescale=1./255,
rotation_range=40,
width_shift_range=0.2,
height_shift_range=0.2,
shear_range=0.2,
zoom_range=0.2,
horizontal_flip=True,)
test_data_gen = ImageDataGenerator(rescale=1./255)
```

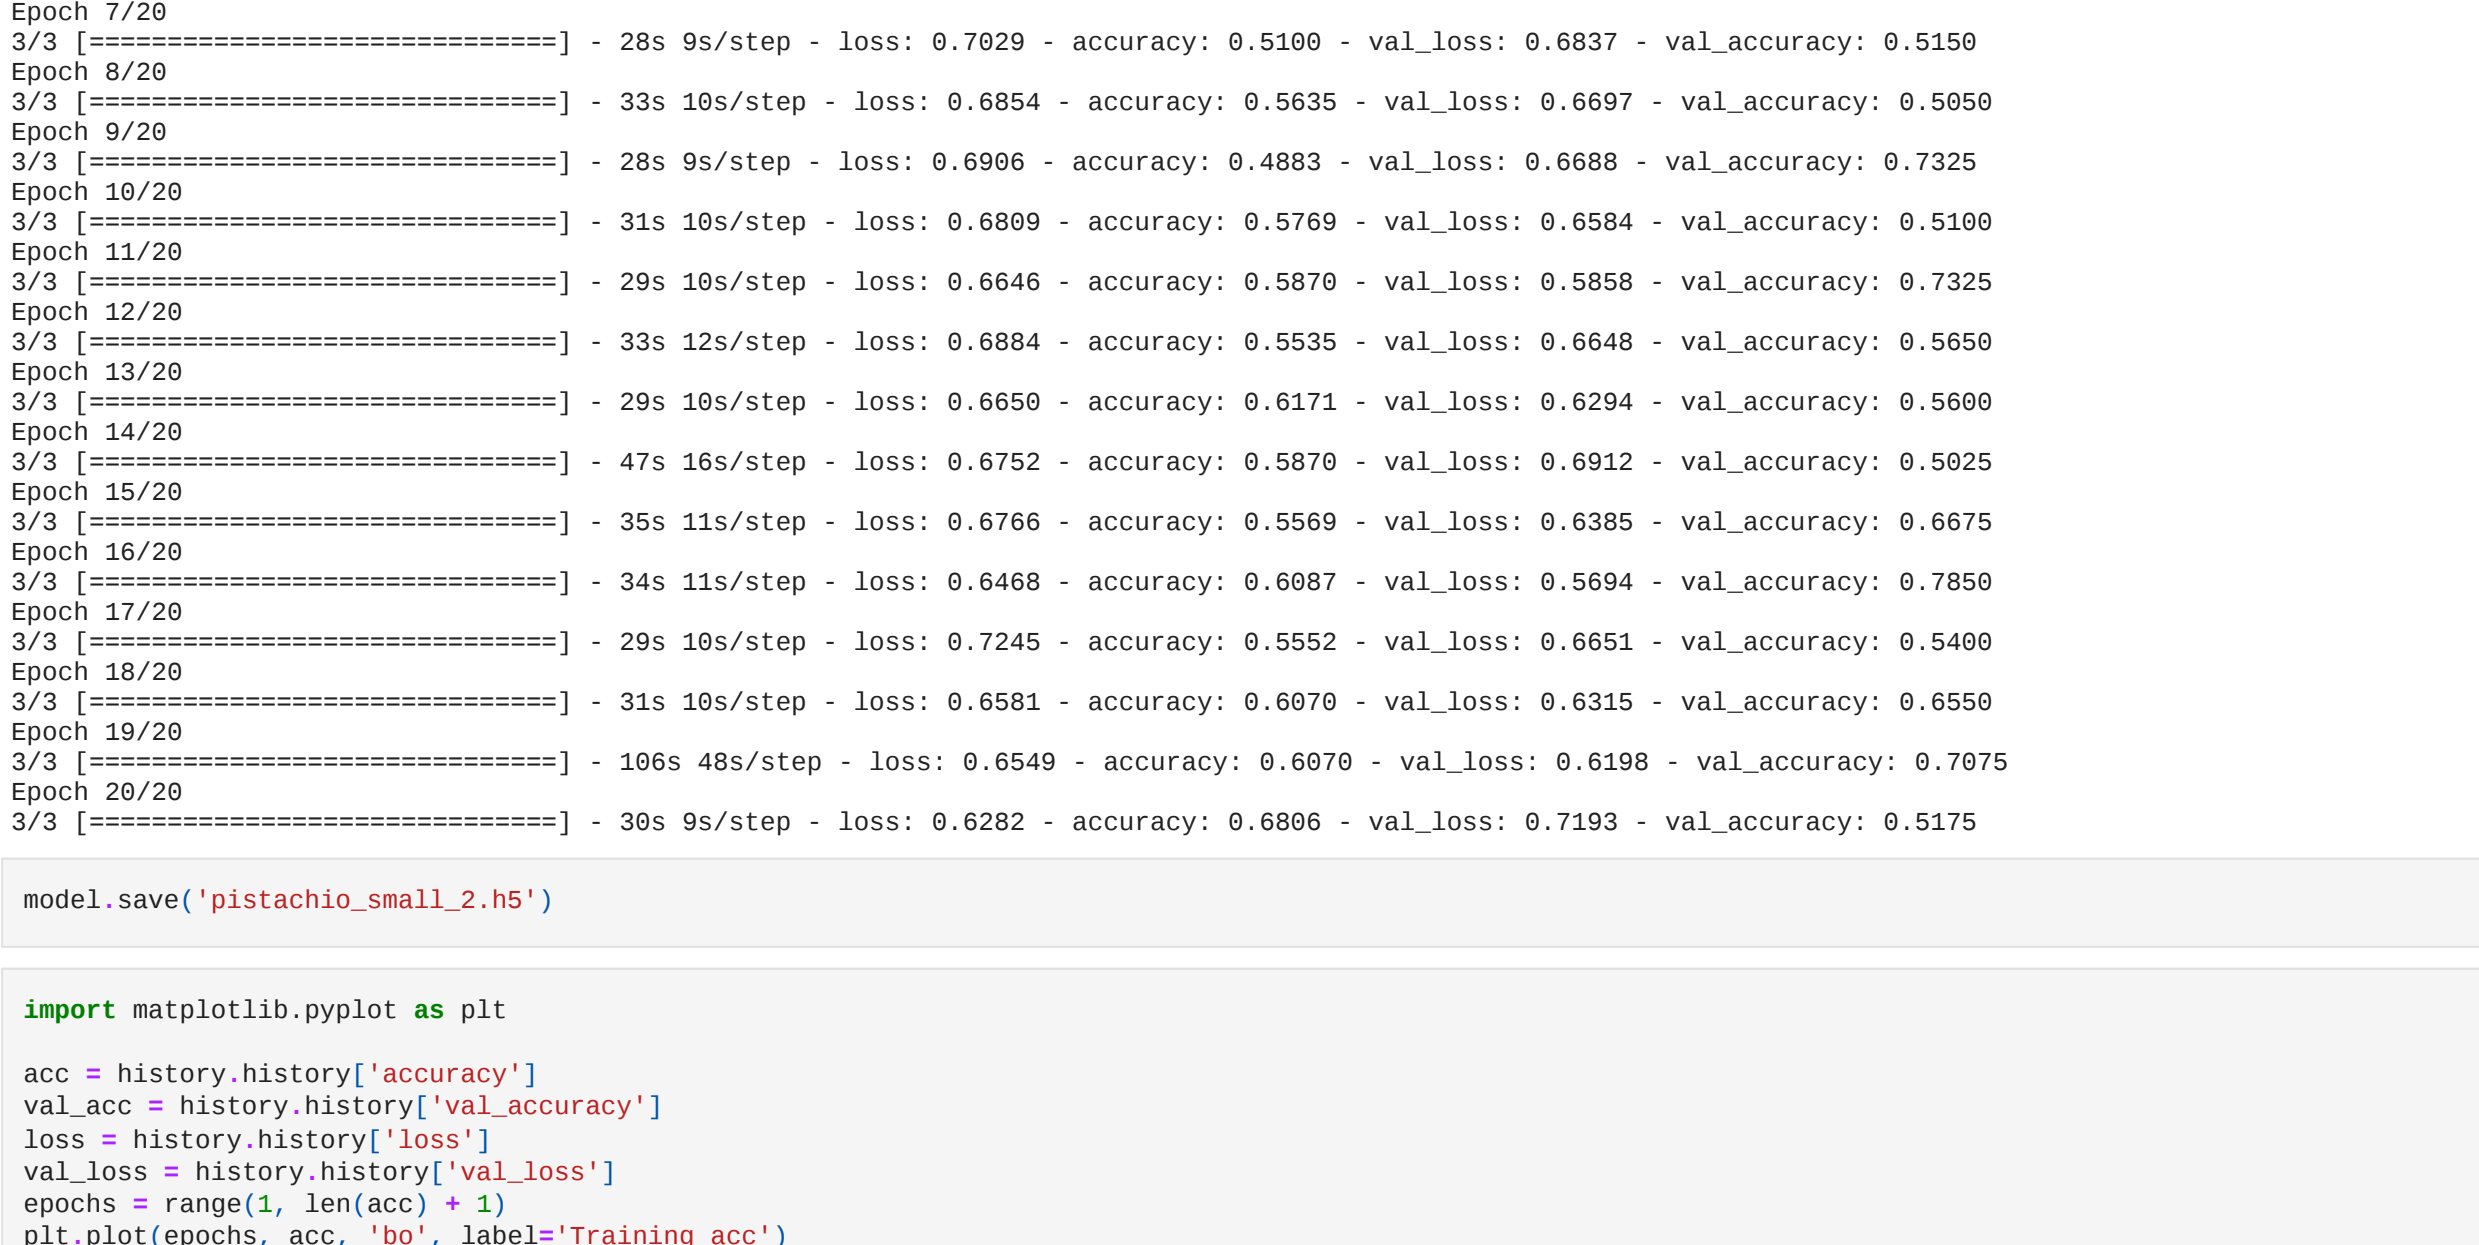
```
train_generator = train_data_gen.flow_from_directory(
    new_base_dir / 'train',
    target_size=(224, 224),
    batch_size=200,
    class_mode='binary',
    shuffle=True)

validation_generator = test_data_gen.flow_from_directory(
    new_base_dir / 'validation',
    target_size=(224, 224),
    batch_size=200,
    class_mode='binary',
    shuffle=True)

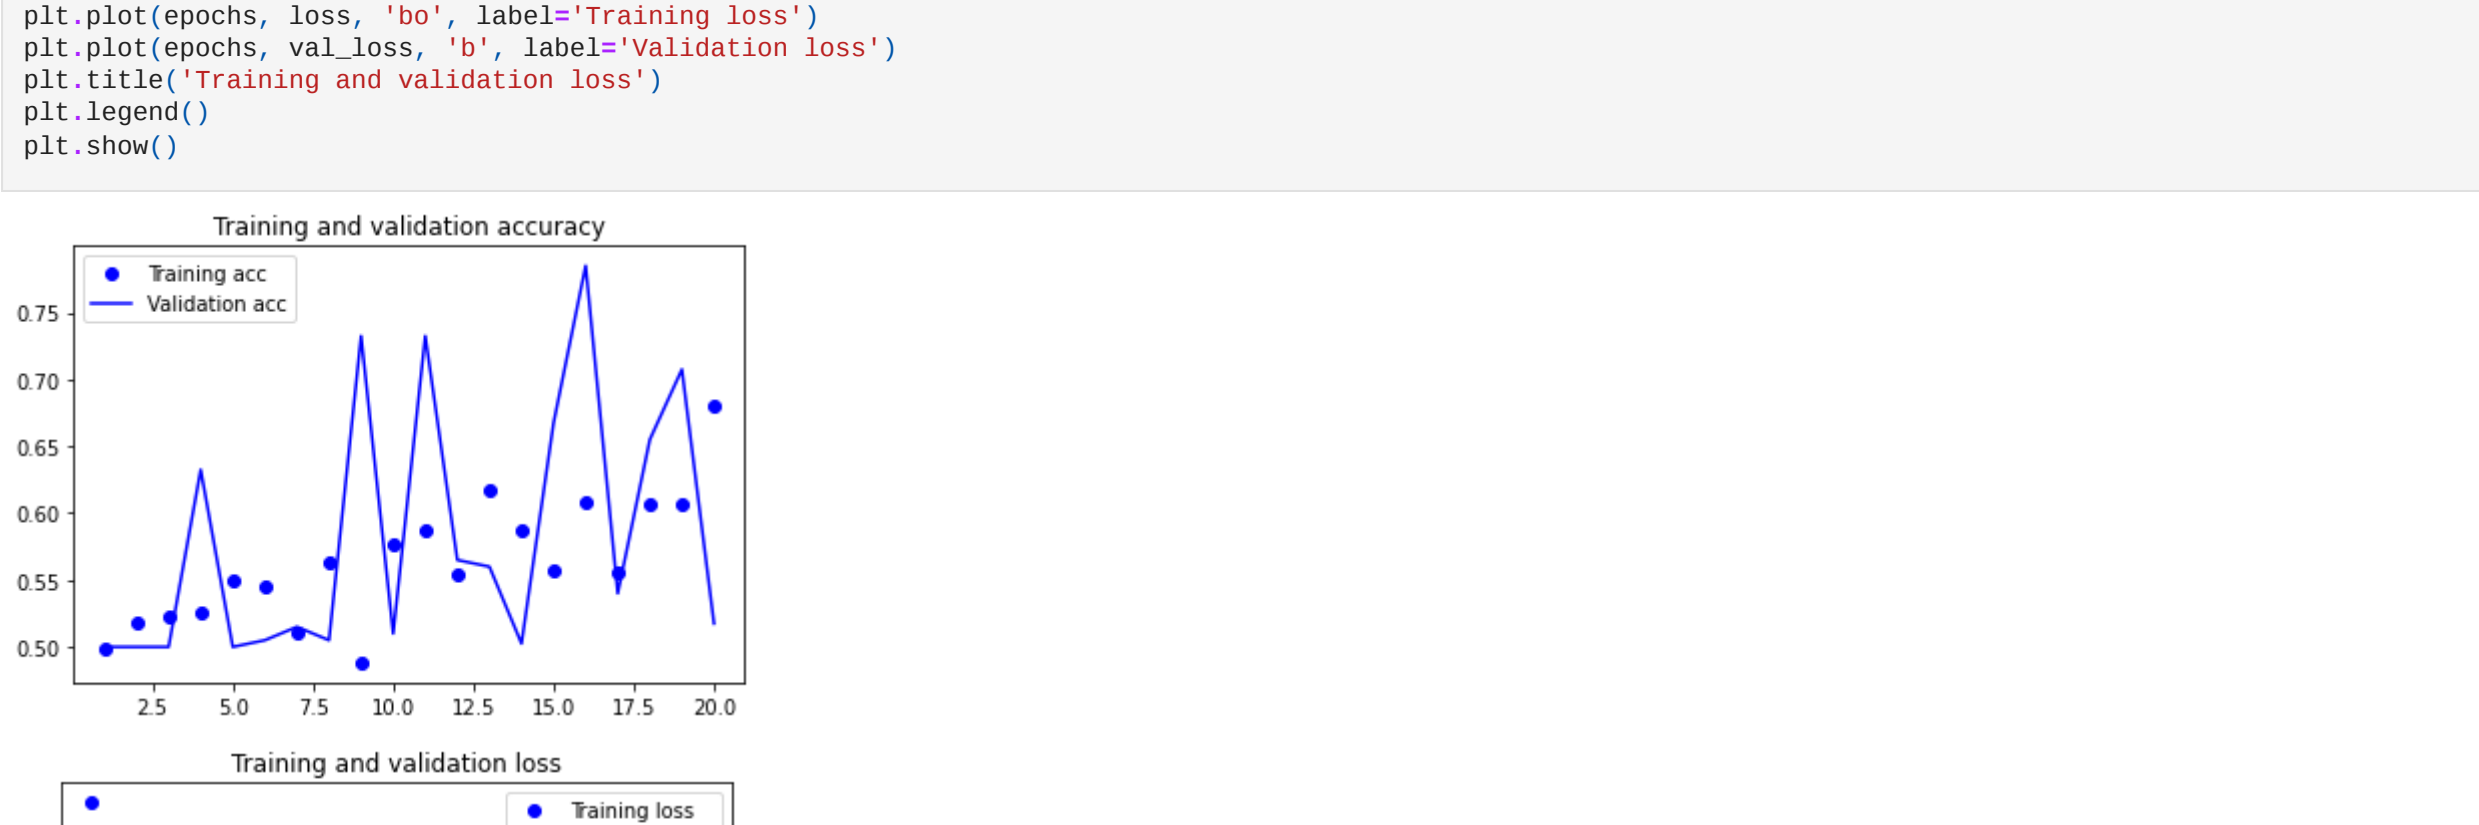
history = model.fit_generator(
    train_generator,
    epochs=20,
    validation_data=validation_generator)
```

Found 598 images belonging to 2 classes.

Found 400 images belonging to 2 classes.



```
In [65]: model.save('pistachio_small_2.h5')
```



In [ ]: